CME 307 / MS&E 311 / OIT 676: Optimization

Optimality conditions and convexity

Professor Udell

Management Science and Engineering
Stanford

October 28, 2024

Outline

Constrained and unconstrained optimization

Optimality conditions

Convex analysis

Convex optimization

Constrained vs unconstrained optimization

constrained optimization

- ▶ examples: scheduling, routing, packing, logistics, scheduling, control
- ▶ what's hard: finding a feasible point

unconstrained optimization

- ▶ examples: data fitting, statistical/machine learning
- ▶ what's hard: reducing the objective

both are necessary for real-world problems!

Unconstrained smooth optimization

for $f : \mathbf{R}^n \to \mathbf{R}$ ctsly differentiable,

$$\begin{array}{ll} \text{minimize} & f(x) \\ \text{variable} & x \in \mathbf{R}^n \end{array}$$

examples:

- ▶ least squares
- ▶ logistic regression
- ▶ neural network training (with smooth activation like tanh, ELU, GeLU, . . .)
- ▶ . . .

Oracles

an optimization **oracle** is your interface for accessing the problem data:
e.g., an oracle for $f : \mathbf{R}^n \to \mathbf{R}$ can evaluate for any $x \in \mathbf{R}^n$:

- **zero-order:** $f_0(x)$
- **first-order:** $f_0(x)$ and $\nabla f_0(x)$
- **second-order:** $f_0(x)$, $\nabla f_0(x)$, and $\nabla^2 f_0(x)$

why oracles?

- can optimize real systems based on observed output (not just models)
- can use and extend old or complex but trusted code (*e.g.*, NASA, PDE simulations, ...)
- can prove lower bounds on the oracle complexity of a problem class

source: Nesterov 2004 "Introductory Lectures on Convex Optimization"

Outline

Solution of an optimization problem

$$\text{minimize} \quad f(x)$$

for $f : \mathcal{D} \to \mathbf{R}$. $x^\star$ is a

▶ **global minimizer** if $f(x) \geq f(x^\star)$ for all $x \in \mathcal{D}$.

▶ **local minimizer** if there is a neighborhood $\mathcal{N}$ around $x^\star$ so that $f(x) \geq f(x^\star)$ for all $x \in \mathcal{N}$.

▶ **isolated local minimizer** if the neighborhood $\mathcal{N}$ contains no other local minimizers.

▶ **unique minimizer** if it is the only global minimizer.

Solution of an optimization problem

$$\text{minimize} \quad f(x)$$

for $f : \mathcal{D} \to \mathbf{R}$. $x^\star$ is a

- **global minimizer** if $f(x) \geq f(x^\star)$ for all $x \in \mathcal{D}$.
- **local minimizer** if there is a neighborhood $\mathcal{N}$ around $x^\star$ so that $f(x) \geq f(x^\star)$ for all $x \in \mathcal{N}$.
- **isolated local minimizer** if the neighborhood $\mathcal{N}$ contains no other local minimizers.
- **unique minimizer** if it is the only global minimizer.

pictures!

First order optimality condition

Theorem

If $x^\star \in \mathbf{R}^n$ is a local minimizer of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$, then $\nabla f(x^\star) = 0$.

First order optimality condition

Theorem

If $x^\star \in \mathbf{R}^n$ is a local minimizer of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$, then $\nabla f(x^\star) = 0$.

proof: suppose by contradiction that $\nabla f(x^\star) \neq 0$. consider points of the form $x_\alpha = x^\star - \alpha \nabla f(x^\star)$ for $\alpha > 0$. by definition of the gradient,

$$\lim_{\alpha \to 0} \frac{f(x_\alpha) - f(x^\star)}{\alpha} = -\nabla f(x^\star)^\top \nabla f(x^\star) = -\|\nabla f(x^\star)\|^2 < 0$$

so for any sufficiently small $\alpha > 0$, we have $f(x_\alpha) < f(x^\star)$, which contradicts the fact that $x^\star$ is a local minimizer.

Second order optimality condition

Theorem

If $x^\star \in \mathbf{R}^n$ is a local minimizer of a twice differentiable function $f : \mathbf{R}^n \to \mathbf{R}$, then $\nabla^2 f(x^\star) \succeq 0$.

Second order optimality condition

Theorem

If $x^\star \in \mathbf{R}^n$ is a local minimizer of a twice differentiable function $f : \mathbf{R}^n \to \mathbf{R}$, then $\nabla^2 f(x^\star) \succeq 0$.

proof: similar to the previous proof. use the fact that the second order approximation

$$f(x_\alpha) \approx f(x^\star) + \nabla f(x^\star)^\top (x_\alpha - x^\star) + \frac{1}{2}(x_\alpha - x^\star)^\top \nabla^2 f(x^\star)(x_\alpha - x^\star)$$

is accurate locally to show a contradiction unless $\nabla^2 f(x^\star) \succeq 0$: if not, there is a direction v such that $v^T \nabla^2 f(x^\star) v < 0$. then $f(x + \alpha v) < f(x^\star)$ for α arbitrarily small, which contradicts the fact that $x^\star$ is a local minimizer.

Symmetric positive semidefinite matrices

Definition

a symmetric matrix $Q \in \mathbf{R}^{n \times n}$ is **positive semidefinite** (psd) if $x^T Q x \geq 0$ for all $x \in \mathbf{R}^n$.

these matrices are so important that there are many ways to write them! for $Q \in \mathbf{R}^{n \times n}$,

$$Q \in \mathbf{S}_+^n \iff Q \succeq 0 \iff Q = Q^T, \ \lambda_{\min}(Q) \geq 0 \iff v^T Q v \geq 0 \quad \forall v \in \mathbf{R}^n$$

Symmetric positive semidefinite matrices

Definition

a symmetric matrix $Q \in \mathbf{R}^{n \times n}$ is **positive semidefinite** (psd) if $x^T Q x \geq 0$ for all $x \in \mathbf{R}^n$.

these matrices are so important that there are many ways to write them! for $Q \in \mathbf{R}^{n \times n}$,

$$Q \in \mathbf{S}_+^n \iff Q \succeq 0 \iff Q = Q^T, \ \lambda_{\min}(Q) \geq 0 \iff v^T Q v \geq 0 \quad \forall v \in \mathbf{R}^n$$

$Q \in \mathbf{S}_+^n$ is **symmetric positive definite** (spd) $(Q \succ 0)$ if $x^T Q x > 0$ for all $x \in \mathbf{R}^n$.

Symmetric positive semidefinite matrices

Definition

a symmetric matrix $Q \in \mathbf{R}^{n \times n}$ is **positive semidefinite** (psd) if $x^T Q x \geq 0$ for all $x \in \mathbf{R}^n$.

these matrices are so important that there are many ways to write them! for $Q \in \mathbf{R}^{n \times n}$,

$$Q \in \mathbf{S}_+^n \iff Q \succeq 0 \iff Q = Q^T, \; \lambda_{\min}(Q) \geq 0 \iff v^T Q v \geq 0 \quad \forall v \in \mathbf{R}^n$$

$Q \in \mathbf{S}_+^n$ is **symmetric positive definite** (spd) ($Q \succ 0$) if $x^T Q x > 0$ for all $x \in \mathbf{R}^n$.

why care about psd matrices Q?

- ▶ least-squares objective has a psd $Q = A^T A$
- ▶ level sets of $x^T Q x$ are (bounded) ellipsoids
- ▶ the quadratic form $x^T Q x$ is a metric iff $Q \succ 0$
- ▶ eigenvalue decomp and svd coincide for psd matrices

Outline

Convex sets

Definition

A set $S \subseteq \mathbf{R}^n$ is convex if it contains every chord: for all $\theta \in [0, 1]$, w, $v \in S$,

$$\theta w + (1 - \theta)v \in S$$

Convex sets

Definition

A set $S \subseteq \mathbf{R}^n$ is convex if it contains every chord: for all $\theta \in [0, 1]$, w, $v \in S$,

$$\theta w + (1 - \theta)v \in S$$

Q: Which of these are convex?
ellipsoid, crescent moon, . . .

Operations that preserve convexity

if S and T are convex, then so are:

- $S \cap T$
- $S + T = \{ s + t \mid s \in S,\ t \in T \}$

Convex functions

a function $f : \mathbf{R}^n \to \mathbf{R}$ is convex iff

Convex functions

a function $f : \mathbf{R}^n \to \mathbf{R}$ is convex iff

▶ **Chords.** it never lies above its chord: $\forall \theta \in [0, 1]$, $w, v \in \mathbf{R}^n$

$$f(\theta w + (1 - \theta)v) \le \theta f(w) + (1 - \theta)f(v)$$

Convex functions

a function $f : \mathbf{R}^n \to \mathbf{R}$ is convex iff

▶ **Chords.** it never lies above its chord: $\forall \theta \in [0, 1]$, $w, v \in \mathbf{R}^n$

$$f(\theta w + (1 - \theta)v) \leq \theta f(w) + (1 - \theta)f(v)$$

▶ **Epigraph.** $\mathbf{epi}(f) = \{(x, t) : t \geq f(x)\}$ is convex

Convex functions

a function $f : \mathbf{R}^n \to \mathbf{R}$ is convex iff

► **Chords.** it never lies above its chord: $\forall \theta \in [0,1]$, $w, v \in \mathbf{R}^n$

$$f(\theta w + (1-\theta)v) \le \theta f(w) + (1-\theta)f(v)$$

► **Epigraph. epi**$(f) = \{(x,t) : t \ge f(x)\}$ is convex
► **First order condition.** if f is differentiable,

$$f(v) \ge f(w) + \nabla f(w)^\top (v - w), \qquad \forall w, v \in \mathbf{R}^n$$

Convex functions

a function $f : \mathbf{R}^n \to \mathbf{R}$ is convex iff

▶ **Chords.** it never lies above its chord: $\forall \theta \in [0, 1]$, $w, v \in \mathbf{R}^n$

$$f(\theta w + (1 - \theta)v) \leq \theta f(w) + (1 - \theta)f(v)$$

▶ **Epigraph. epi**$(f) = \{(x, t) : t \geq f(x)\}$ is convex

▶ **First order condition.** if f is differentiable,

$$f(v) \geq f(w) + \nabla f(w)^\top (v - w), \qquad \forall w, v \in \mathbf{R}^n$$

▶ **Second order condition.** If f is twice differentiable, its Hessian is always psd:

$$\lambda_{\min}(\nabla^2 f(x)) \geq 0, \qquad \forall x \in \mathbf{R}^n$$

Convexity examples

Q: Which of these functions are convex?

- quadratic function $f(x) = x^2$ for $x \in \mathbf{R}$
- absolute value function $f(x) = |x|$ for $x \in \mathbf{R}$
- quadratic function $f(x) = x^T A x$, $x \in \mathbf{R}^n$, $A \succeq 0$
- quadratic function $f(x) = x^T A x$, A indefinite
- rollercoaster function (cubic) $f(x) = (x-1)(x-3)(x-5)$
- hyperbolic function $f(x) = 1/x$ for $x > 0$
- jump function $f(x) = 1$ if $x \geq 0$, $f(x) = 0$ otherwise
- jump to infinity function $f(x) = 1$ if $x \in [-1, 1]$, $f(x) = \infty$ otherwise

Operations that preserve convexity

if $f : \mathbf{R}^n \to \mathbf{R}$ and $g : \mathbf{R}^n \to \mathbf{R}$ are convex, then so are:

- cf for $c \geq 0$
- $f(Ax + b)$ for $A \in \mathbf{R}^n \times m$, $b \in \mathbf{R}^n$
- $f + g$
- $\max\{f, g\}$

Operations that preserve convexity

if $f : \mathbf{R}^n \to \mathbf{R}$ and $g : \mathbf{R}^n \to \mathbf{R}$ are convex, then so are:

- cf for $c \geq 0$
- $f(Ax + b)$ for $A \in \mathbf{R}^n \times m$, $b \in \mathbf{R}^n$
- $f + g$
- $\max\{f, g\}$

Q: Pick one and assume f and g are twice-differentiable. What is the easiest way to prove convexity?

Operations that preserve convexity

if $f : \mathbf{R}^n \to \mathbf{R}$ and $g : \mathbf{R}^n \to \mathbf{R}$ are convex, then so are:

- cf for $c \geq 0$
- $f(Ax + b)$ for $A \in \mathbf{R}^n \times m$, $b \in \mathbf{R}^n$
- $f + g$
- $\max\{f, g\}$

Q: Pick one and assume f and g are twice-differentiable. What is the easiest way to prove convexity?

most general rule:

$f \circ g(x) = f(g(x))$ is convex if g is convex and f is convex and nondecreasing

since

$$(f \circ g)''(x) = f''(g(x))(g'(x))^2 + f'(g(x))g''(x)$$

Jensen's inequality

Jensen's inequality generalizes the first-order condition to distribution of points:

Theorem

If $f : \mathbf{R}^n \to \mathbf{R}$ is convex and X is a random variable, then

$$f(\mathbb{E}[X]) \leq \mathbb{E}[f(X)]$$

Sublevel set

Definition

The **sublevel set** of a function $f : \mathbf{R}^n \to \mathbf{R}$ at level t is

$$S_t = \{x \in \mathbf{R}^n \mid f(x) \leq t\}$$

Sublevel set

Definition

The **sublevel set** of a function $f : \mathbf{R}^n \to \mathbf{R}$ at level t is

$$S_t = \{x \in \mathbf{R}^n \mid f(x) \leq t\}$$

Theorem

A convex function $f : \mathbf{R}^n \to \mathbf{R}$ has convex sublevel sets.

Sublevel set

Definition

The **sublevel set** of a function $f : \mathbf{R}^n \to \mathbf{R}$ at level t is

$$S_t = \{x \in \mathbf{R}^n \mid f(x) \le t\}$$

Theorem

A convex function $f : \mathbf{R}^n \to \mathbf{R}$ has convex sublevel sets.

proof: Jensen's inequality. if $x, y \in S_t$, then for $\theta \in [0, 1]$,

$$f(\theta x + (1 - \theta)y) \le \theta f(x) + (1 - \theta)f(y) \le \theta t + (1 - \theta)t = t$$

so $\theta x + (1 - \theta)y \in S_t$.

Quasiconvexity

converse is not true: a function can have all sublevel sets convex, and still be non-convex.

Definition

A function $f : \mathbf{R}^n \to \mathbf{R}$ is **quasiconvex** if its sublevel sets are convex.

examples of functions that are quasiconvex but not convex?

Supporting hyperplane

Definition

A **supporting hyperplane** to a set $S \subseteq \mathbf{R}^n$ at a point $x \in S$ is a hyperplane that touches S at x and lies entirely on one side of S:

$$H = \{y \in \mathbf{R}^n \mid a^\top y = b\} \text{ supports } S \text{ at } x \text{ if} \qquad \begin{aligned} a^\top x &= b \\ a^\top y &\geq b \quad \forall y \in S \end{aligned}$$

Supporting hyperplane

Definition

A **supporting hyperplane** to a set $S \subseteq \mathbf{R}^n$ at a point $x \in S$ is a hyperplane that touches S at x and lies entirely on one side of S:

$$H = \{y \in \mathbf{R}^n \mid a^\top y = b\} \text{ supports } S \text{ at } x \text{ if} \qquad \begin{aligned} a^\top x &= b \\ a^\top y &\geq b \quad \forall y \in S \end{aligned}$$

Theorem (Supporting hyperplane)

Any nonempty convex set has a supporting hyperplane at every boundary point.

Supporting hyperplane condition for convexity

Theorem (Partial converse)

If a closed set with nonempty interior has a supporting hyperplane at every boundary point, then it is convex.

Supporting hyperplane condition for convexity

Theorem (Partial converse)

If a closed set with nonempty interior has a supporting hyperplane at every boundary point, then it is convex.

Theorem

A function $f : \mathbf{R}^n \to \mathbf{R}$ is convex $\iff$ for all $x \in \mathbf{dom}\, f$, the epigraph of f has a supporting hyperplane at $(x, f(x))$: for some $g \in \mathbf{R}^n$,

$$f(y) \geq f(x) + g^\top (y - x) \quad \forall y \in \mathbf{R}^n$$

Supporting hyperplane condition for convexity

Theorem (Partial converse)

If a closed set with nonempty interior has a supporting hyperplane at every boundary point, then it is convex.

Theorem

A function $f : \mathbf{R}^n \to \mathbf{R}$ is convex $\iff$ for all $x \in \mathbf{dom}\, f$, the epigraph of f has a supporting hyperplane at $(x, f(x))$: for some $g \in \mathbf{R}^n$,

$$f(y) \geq f(x) + g^\top(y - x) \quad \forall y \in \mathbf{R}^n$$

generalizes first-order condition for convexity to non-differentiable functions!

Supporting hyperplane condition for convexity

Theorem (Partial converse)

If a closed set with nonempty interior has a supporting hyperplane at every boundary point, then it is convex.

Theorem

A function $f : \mathbf{R}^n \to \mathbf{R}$ is convex $\iff$ for all $x \in \mathbf{dom}\, f$, the epigraph of f has a supporting hyperplane at $(x, f(x))$: for some $g \in \mathbf{R}^n$,

$$f(y) \geq f(x) + g^\top (y - x) \quad \forall y \in \mathbf{R}^n$$

generalizes first-order condition for convexity to non-differentiable functions!

Definition

A vector $g \in \mathbf{R}^n$ is a **subgradient** of $f : \mathbf{R}^n \to \mathbf{R}$ at $x \in \mathbf{R}^n$ if
$f(y) \geq f(x) + g^\top (y - x)$ for all $y \in \mathbf{R}^n$.

$f = \max\{f_1, f_2\}$, with f_1, f_2 convex and differentiable

Q: Where is the function f differentiable? Where is the subgradient unique?

Subdifferential

set of all subgradients of f at x is called the **subdifferential** $\partial f(x)$

$$\partial f(x) = \{g : f(y) \geq f(x) + g^T(y - x) \quad \forall y\}$$

Subdifferential

set of all subgradients of f at x is called the **subdifferential** $\partial f(x)$

$$\partial f(x) = \{g : f(y) \geq f(x) + g^T(y - x) \quad \forall y\}$$

for any f,

▶ $\partial f(x)$ is a closed convex set (can be empty)
▶ $\partial f(x) = \emptyset$ if $f(x) = \infty$

proof: use the definition

Subdifferential

set of all subgradients of f at x is called the **subdifferential** $\partial f(x)$

$$\partial f(x) = \{g : f(y) \geq f(x) + g^T(y - x) \quad \forall y\}$$

for any f,

- ▶ $\partial f(x)$ is a closed convex set (can be empty)
- ▶ $\partial f(x) = \emptyset$ if $f(x) = \infty$

proof: use the definition

if f is convex,

- ▶ $\partial f(x)$ is nonempty, for $x \in$ **relint dom** f
- ▶ $\partial f(x) = \{\nabla f(x)\}$, if f is differentiable at x
- ▶ if $\partial f(x) = \{g\}$, then f is differentiable at x and $g = \nabla f(x)$

Outline

Convex optimization

an optimization problem is convex if:

- **Geometrically:** the feasible set and the epigraph of the objective are convex
- **NLP:** the objective and inequality constraints are convex functions, and the equality constraints are affine

Convex optimization

an optimization problem is convex if:

- ▶ **Geometrically:** the feasible set and the epigraph of the objective are convex
- ▶ **NLP:** the objective and inequality constraints are convex functions, and the equality constraints are affine

why convex optimization?

- ▶ relatively complete theory
- ▶ efficient solvers
- ▶ conceptual tools that generalize linear programming: duality, stopping conditions, . . .

Convex optimization

an optimization problem is convex if:

▶ **Geometrically:** the feasible set and the epigraph of the objective are convex
▶ **NLP:** the objective and inequality constraints are convex functions, and the equality constraints are affine

why convex optimization?

▶ relatively complete theory
▶ efficient solvers
▶ conceptual tools that generalize linear programming: duality, stopping conditions, . . .

▶ a function f is concave if $-f$ is convex
▶ concave maximization $\implies$ a convex optimization problem

Local minima are global for convex functions

Theorem

If $x^\star$ is a local minimizer of a convex function f, then $x^\star$ is a global minimizer.

Local minima are global for convex functions

Theorem

If $x^\star$ is a local minimizer of a convex function f, then $x^\star$ is a global minimizer.

proof?

Local minima are global for convex functions

Theorem

If $x^\star$ is a local minimizer of a convex function f, then $x^\star$ is a global minimizer.

proof? suppose by contradiction that another point x' is a global minimizer, with $f(x') < f(x^\star)$. draw the chord between x' and $x^\star$. since the chord lies above f, every convex combination $x = \theta x^\star + (1 - \theta)x'$ of x' and $x^\star$ for $\theta \in (0, 1)$ has a value $f(x) < f(x^\star)$. this is true even for $x \to x^\star$, contradicting our assumption that $x^\star$ is a local minimizer.

Corollary

Corollary

If f is convex and differentiable and $\nabla f(x^\star) = 0$, then $x^\star$ is a global minimizer.

Corollary

If f is convex and differentiable and $\nabla f(x^\star) = 0$, then $x^\star$ is a global minimizer.

Q: Is a global minimizer of a convex function always unique?

Corollary

Corollary

If f is convex and differentiable and $\nabla f(x^\star) = 0$, then $x^\star$ is a global minimizer.

Q: Is a global minimizer of a convex function always unique?
A: No. Picture.

First-order condition

Definition

$x^\star \in \mathbf{R}^n$ is a **stationary point** of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ if $\nabla f(x^\star) = 0$.

First-order condition

Definition

$x^\star \in \mathbf{R}^n$ is a **stationary point** of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ if $\nabla f(x^\star) = 0$.

Q: Can a global minimum have a non-zero gradient?

First-order condition

Definition

$x^\star \in \mathbf{R}^n$ is a **stationary point** of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ if $\nabla f(x^\star) = 0$.

Q: Can a global minimum have a non-zero gradient?
A: No.

First-order condition

Definition

$x^\star \in \mathbf{R}^n$ is a **stationary point** of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ if $\nabla f(x^\star) = 0$.

Q: Can a global minimum have a non-zero gradient?
A: No.
Q: Is a stationary point always a global minimum?

First-order condition

Definition

$x^\star \in \mathbf{R}^n$ is a **stationary point** of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ if $\nabla f(x^\star) = 0$.

Q: Can a global minimum have a non-zero gradient?
A: No.
Q: Is a stationary point always a global minimum?
A: No.

First-order condition

Definition

$x^\star \in \mathbf{R}^n$ is a **stationary point** of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ if $\nabla f(x^\star) = 0$.

Q: Can a global minimum have a non-zero gradient?
A: No.
Q: Is a stationary point always a global minimum?
A: No.
Q: ... for convex functions?

First-order condition

Definition

$x^\star \in \mathbf{R}^n$ is a **stationary point** of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ if $\nabla f(x^\star) = 0$.

Q: Can a global minimum have a non-zero gradient?
A: No.
Q: Is a stationary point always a global minimum?
A: No.
Q: . . . for convex functions?
A: Yes.

First-order condition

Definition

$x^\star \in \mathbf{R}^n$ is a **stationary point** of a differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ if $\nabla f(x^\star) = 0$.

Q: Can a global minimum have a non-zero gradient?

A: No.

Q: Is a stationary point always a global minimum?

A: No.

Q: . . . for convex functions?

A: Yes.

$\nabla f(x^\star) = 0$ is the **first-order (necessary) condition** for optimality.

Invex function

Definition

A differentiable function $f : \mathbf{R}^n \to \mathbf{R}$ is **invex** if for some vector-valued function $\eta : \mathbf{R}^n \times \mathbf{R}^n \to \mathbf{R}^n$,

$$f(x) - f(u) \geq \eta(x, u)^\top \nabla f(u) \qquad \forall u \in \mathbf{R}^n, \ x \in \mathbf{dom} \, f$$

Theorem (Craven and Glover, Ben-Israel and Mond)

A function is invex iff every stationary point is a global minimum.

why invex?

▶ generalizes convexity

▶ broadest class of functions for which every stationary point is a global minimum