

# CME 307 / MS&E 311: Optimization

## Least squares

Professor Udell

Management Science and Engineering  
Stanford

April 17, 2023

## Announcements

- ▶ 1:30pm Friday 4/14: team formation in Thornton 110
- ▶ homework 1 out, due Friday 4/21

## Linear system

find  $x \in \mathbf{R}^n$  such that

$$Ax = b$$

given **design matrix**  $A \in \mathbf{R}^{m \times n}$ , **righthand side** (rhs)  $b \in \mathbf{R}^m$

how to solve?

- ▶ factor and solve
  - ▶ QR
  - ▶ singular value decomposition (SVD)
  - ▶ Cholesky (for symmetric  $A$ )
- ▶ iterative methods
  - ▶ conjugate gradient (CG) (for symmetric  $A$ )
  - ▶ iterative refinement

we will talk about QR, CG, and iterative refinement

## The SVD and the pseudoinverse

if  $r = \text{Rank}(A)$  and  $A = U\Sigma V^T$  is the SVD of  $A$ ,

- ▶  $U \in \mathbf{R}^{m \times r}$  is orthogonal:  $U^T U = I_r$
- ▶  $\Sigma \in \mathbf{R}_+^{r \times r}$  is diagonal and nonnegative
- ▶  $V \in \mathbf{R}^{n \times r}$  is orthogonal:  $V^T V = I_r$

we can write the **pseudoinverse**  $A^\dagger = V\Sigma^{-1}U^T$

- ▶ if  $x \in \text{span}(V)$ ,  $A^\dagger A x = x$

## Considerations in choosing a method

- ▶ sparse or dense  $A$ ?
- ▶ symmetric  $A$  or rectangular problem?
- ▶ conditioning of  $A$ ?
- ▶ one problem, or many righthand sides  $b$  with the same design matrix  $A$ ?

	<b>symmetric psd</b>	<b>rectangular</b>
<b>direct</b>	Cholesky	QR
<b>indirect</b>	CG	LSQR

Table: Methods for solving linear systems

- ▶ direct methods get accurate solutions in  $O(n^3)$  flops
- ▶ indirect methods get ok solution in a small number of matvecs

## Optimality condition for least squares is a linear system

given  $A \in \mathbf{R}^{m \times n}$ ,  $y \in \mathbf{R}^m$ . find  $x$  to solve

$$\text{minimize} \quad \|Ax - b\|^2.$$

to solve, take gradient, set to 0. solution  $x$  satisfies **normal equations**

$$A^\top Ax = A^\top b.$$

a linear system!

- ▶  $A^\top A$  symmetric positive semidefinite
- ▶ normal equations always have a solution (why?)

# Outline

QR

Conjugate gradient

Preconditioned CG

Iterative refinement

## How to solve a linear system?

**never form the inverse explicitly:** numerically unstable!

Corollary: never type `inv(A'*A)` or `pinv(A'*A)` to solve the normal equations.



## How to solve a linear system?

**never form the inverse explicitly:** numerically unstable!

Corollary: never type `inv(A'*A)` or `pinv(A'*A)` to solve the normal equations.

Instead: compute the inverse using easier matrices to invert, like

- ▶ orthogonal matrix  $Q$ :

$$a = Qb \iff Q^T a = b$$

- ▶ (upper) triangular matrix  $R$ :  
if  $a = Rb$ , can find  $b$  given  $R$  and  $a$  by solving sequence of simple, stable equations.

## The QR factorization

every matrix  $A$  can be written using **QR decomposition** as  $A = QR$

- ▶  $Q \in \mathbf{R}^{m \times n}$  has orthogonal columns:  $Q^\top Q = I_n$
- ▶  $R \in \mathbf{R}^{n \times n}$  is upper triangular:  $R_{ij} = 0$  for  $i > j$
- ▶ diagonal of  $R \in \mathbf{R}^{n \times n}$  is positive:  $R_{ii} > 0$  for  $i = 1, \dots, n$
- ▶ this factorization always exists and is unique  
(proof by Gram-Schmidt construction)

can compute  $QR$  factorization of  $X$  in  $2mn^2$  flops

## The QR factorization

every matrix  $A$  can be written using **QR decomposition** as  
 $A = QR$

- ▶  $Q \in \mathbf{R}^{m \times n}$  has orthogonal columns:  $Q^T Q = I_n$
- ▶  $R \in \mathbf{R}^{n \times n}$  is upper triangular:  $R_{ij} = 0$  for  $i > j$
- ▶ diagonal of  $R \in \mathbf{R}^{n \times n}$  is positive:  $R_{ii} > 0$  for  $i = 1, \dots, n$
- ▶ this factorization always exists and is unique  
(proof by Gram-Schmidt construction)

can compute QR factorization of  $X$  in  $2mn^2$  flops

use `LinearAlgebra.qr`:

$$Q, R = \mathbf{qr}(X)$$

**advantage of QR:** it's easy to invert  $R$ !

## QR to solve linear system

use QR to solve linear system  $Ax = b$ : if  $A = QR$ ,

$$Ax = b$$

$$x = R^{-1}Q^T b$$

## QR to solve linear system

use QR to solve linear system  $Ax = b$ : if  $A = QR$ ,

$$\begin{aligned} Ax &= b \\ x &= R^{-1}Q^T b \end{aligned}$$

**Q:** What happens if we apply this method to solve an infeasible system with  $m > n$ ?

## QR to solve linear system

use QR to solve linear system  $Ax = b$ : if  $A = QR$ ,

$$\begin{aligned} Ax &= b \\ x &= R^{-1}Q^T b \end{aligned}$$

**Q:** What happens if we apply this method to solve an infeasible system with  $m > n$ ?

**A:** decompose  $b = b^{\parallel} + b^{\perp}$  where  $b^{\parallel} \in \text{span}(A)$ ; QR solves  $Ax = b^{\parallel}$

## QR for least squares

use QR to solve least squares: if  $A = QR$ ,

$$A^{\top}Ax = A^{\top}b$$

$$(QR)^{\top}QRx = (QR)^{\top}b$$

$$R^{\top}Q^{\top}QRx = R^{\top}Q^{\top}b$$

$$R^{\top}Rx = R^{\top}Q^{\top}b$$

$$Rx = Q^{\top}b$$

$$x = R^{-1}Q^{\top}b$$

## Computational considerations

use  $QR$  factorization to solve  $Ax = b$

- ▶ compute  $QR$  factorization of  $A$  ( $2mn^2$  flops)
- ▶ to compute  $x = R^{-1}Q^T b$ 
  - ▶ form  $z = Q^T b$  ( $2mn$  flops)
  - ▶ compute  $x = R^{-1}z$  by back-substitution ( $n^2$  flops)



## Computational considerations

use  $QR$  factorization to solve  $Ax = b$

- ▶ compute  $QR$  factorization of  $A$  ( $2mn^2$  flops)
- ▶ to compute  $x = R^{-1}Q^T b$ 
  - ▶ form  $z = Q^T b$  ( $2mn$  flops)
  - ▶ compute  $x = R^{-1}z$  by back-substitution ( $n^2$  flops)

in julia (or matlab), the **backslash operator** solves least-squares efficiently (usually, using QR)

$$x = A \setminus b$$

in python, use `numpy.linalg.lstsq`

## Demo: QR

[https:  
//github.com/stanford-cme-307/demos/blob/main/lq.ipynb](https://github.com/stanford-cme-307/demos/blob/main/lq.ipynb)

## Sparse QR

complexity of QR depends on the sparsity of  $Q$  and  $R$ :

- ▶ compute  $QR$  factorization of  $A$  (?? flops)
- ▶ to compute  $x = R^{-1}Q^{\top}b$ 
  - ▶ form  $z = Q^{\top}b$  (**nnz**( $Q$ ) flops)
  - ▶ compute  $x = R^{-1}z$  by back-substitution (**nnz**( $R$ ) flops)

## Q-less QR

during QR, can compute  $Q^T b$  essentially for free!

- ▶ compute QR of  $\begin{bmatrix} A & b \end{bmatrix}$ .

## Q-less QR

during QR, can compute  $Q^\top b$  essentially for free!

- compute QR of  $[A \ b]$ .

or compute it afterwards without forming  $Q$ :

$$\begin{aligned} A^\top b &= (QR)^\top b = R^\top Q^\top b \\ R^{-1} A^\top b &= Q^\top b \end{aligned}$$

## Cholesky and QR

consider **Gram matrix**  $G = A^T A \succeq 0$ . if  $A = QR$ ,

$$G = R^T Q^T Q R = R^T R$$

this construction gives **Cholesky factorization** of a spd matrix  $G$

- ▶ factors spd matrix into triangular matrices
- ▶ Cholesky factors of  $X^T X$  have same structure as  $R$

## Sparse QR: exercise

- ▶ can you guess the sparsity of  $R$  given sparsity of  $A$ ?
- ▶ can you change sparsity of  $R$  by permuting columns of  $A$ ?

## Sparse QR: exercise

- ▶ can you guess the sparsity of  $R$  given sparsity of  $A$ ?
- ▶ can you change sparsity of  $R$  by permuting columns of  $A$ ?

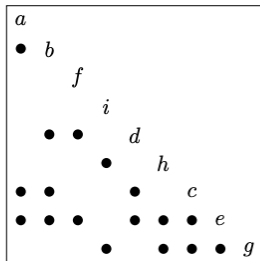

use 'colamd' in Matlab, equivalents in Python and Julia



## Chordal fill-in

to analyze fill-in

- ▶ consider spd matrix, for simplicity
- ▶ interpret matrix as directed graph
- ▶ form clique tree
- ▶ identify fill-in



**Figure 4.1:** *Left.* Filled graph with 9 vertices. The number next to each vertex is the index  $\sigma^{-1}(v)$ . *Right.* Array representation of the same graph.

# Outline

QR

Conjugate gradient

Preconditioned CG

Iterative refinement

## Conjugate gradients

symmetric positive semidefinite system of equations

$$Ax = b, \quad A \in \mathbf{R}^{n \times n}, \quad A = A^T \succeq 0$$

## Conjugate gradients

symmetric positive semidefinite system of equations

$$Ax = b, \quad A \in \mathbf{R}^{n \times n}, \quad A = A^T \succeq 0$$

why use conjugate gradients?

- ▶ uses only matrix-vector multiplies with  $A$ 
  - ▶ useful for structured (from PDE or graph) or sparse matrices, easy to parallelize, ...
- ▶ most useful for problems with  $n > 10^5$  or more
- ▶ converges exactly in  $n$  iterations
- ▶ converges approximately much faster
- ▶ quick-and-dirty solve is appropriate **inside** inner loop of optimization algo

## Conjugate gradients

symmetric positive semidefinite system of equations

$$Ax = b, \quad A \in \mathbf{R}^{n \times n}, \quad A = A^T \succeq 0$$

why use conjugate gradients?

- ▶ uses only matrix-vector multiplies with  $A$ 
  - ▶ useful for structured (from PDE or graph) or sparse matrices, easy to parallelize, ...
- ▶ most useful for problems with  $n > 10^5$  or more
- ▶ converges exactly in  $n$  iterations
- ▶ converges approximately much faster
- ▶ quick-and-dirty solve is appropriate **inside** inner loop of optimization algo

other variants for indefinite (MINRES) or nonsymmetric matrices (GMRES)

## Iterative methods for least squares

define

- ▶ (convex) objective  $f(x) = (1/2)x^\top Ax - x^\top b$
- ▶ gradient  $\nabla f(x) = Ax - b$
- ▶ condition number  $\kappa(A) = \lambda_1(A)/\lambda_n(A)$
- ▶  $A$ -norm  $\|x\|_A^2 = x^\top Ax$
- ▶ bound  $R \geq \|x_\star\|$  on norm of solution  $x_\star$
- ▶ goal: find apx solution within accuracy  $f(x) - f(x_\star) \leq \epsilon$

## Iterative methods for least squares

define

- ▶ (convex) objective  $f(x) = (1/2)x^\top Ax - x^\top b$
- ▶ gradient  $\nabla f(x) = Ax - b$
- ▶ condition number  $\kappa(A) = \lambda_1(A)/\lambda_n(A)$
- ▶  $A$ -norm  $\|x\|_A^2 = x^\top Ax$
- ▶ bound  $R \geq \|x_\star\|$  on norm of solution  $x_\star$
- ▶ goal: find apx solution within accuracy  $f(x) - f(x_\star) \leq \epsilon$

how many iterations (matvecs) required?

## Iterative methods for least squares

define

- ▶ (convex) objective  $f(x) = (1/2)x^\top Ax - x^\top b$
- ▶ gradient  $\nabla f(x) = Ax - b$
- ▶ condition number  $\kappa(A) = \lambda_1(A)/\lambda_n(A)$
- ▶  $A$ -norm  $\|x\|_A^2 = x^\top Ax$
- ▶ bound  $R \geq \|x_\star\|$  on norm of solution  $x_\star$
- ▶ goal: find apx solution within accuracy  $f(x) - f(x_\star) \leq \epsilon$

how many iterations (matvecs) required?

- ▶ conjugate gradient
  - ▶  $O(\sqrt{\kappa} \log(\frac{1}{\epsilon}))$



## Iterative methods for least squares

define

- ▶ (convex) objective  $f(x) = (1/2)x^\top Ax - x^\top b$
- ▶ gradient  $\nabla f(x) = Ax - b$
- ▶ condition number  $\kappa(A) = \lambda_1(A)/\lambda_n(A)$
- ▶  $A$ -norm  $\|x\|_A^2 = x^\top Ax$
- ▶ bound  $R \geq \|x_\star\|$  on norm of solution  $x_\star$
- ▶ goal: find apx solution within accuracy  $f(x) - f(x_\star) \leq \epsilon$

how many iterations (matvecs) required?

- ▶ conjugate gradient
  - ▶  $O(\sqrt{\kappa} \log(\frac{1}{\epsilon}))$
- ▶ gradient descent (GD)
  - ▶  $O(\kappa \log(1/\epsilon))$

## Iterative methods for least squares

define

- ▶ (convex) objective  $f(x) = (1/2)x^\top Ax - x^\top b$
- ▶ gradient  $\nabla f(x) = Ax - b$
- ▶ condition number  $\kappa(A) = \lambda_1(A)/\lambda_n(A)$
- ▶ A-norm  $\|x\|_A^2 = x^\top Ax$
- ▶ bound  $R \geq \|x_\star\|$  on norm of solution  $x_\star$
- ▶ goal: find apx solution within accuracy  $f(x) - f(x_\star) \leq \epsilon$

how many iterations (matvecs) required?

- ▶ conjugate gradient
  - ▶  $O(\sqrt{\kappa} \log(\frac{1}{\epsilon}))$
- ▶ gradient descent (GD)
  - ▶  $O(\kappa \log(1/\epsilon))$
- ▶ accelerated gradient descent
  - ▶  $O(\sqrt{\kappa} \log(\frac{R^2}{\epsilon}))$  more generalizable, but more parameters to tune

source: Bubeck, 2014; Karimi, Nutini, and Schmidt, 2016

## Residual

define **residual**  $r = b - Ax$  at putative solution  $x$

►  $r = -\nabla f(x) = A(x_\star - x)$

## Residual

define **residual**  $r = b - Ax$  at putative solution  $x$

▶  $r = -\nabla f(x) = A(x_\star - x)$

measures of error:

- ▶ objective function  $f(x) - f(x_\star)$
- ▶ norm of residual  $\|r\|$
- ▶ norm of gradient  $\|\nabla f(x)\|$
- ▶ in terms of  $r$ , can compute error in objective

$$\begin{aligned} f(x) - f(x_\star) &= \|x - x_\star\|_A \\ &= \frac{1}{2}(x - x_\star)^\top A(x - x_\star) \\ &= \frac{1}{2}(r)^\top A^{-1}(r) \\ &= \|r\|_{A^{-1}} \end{aligned}$$

## Krylov subspace

the Krylov subspace of dimension  $k$  is

$$\mathcal{K}_k = \text{span}\{b, Ab, \dots, A^{k-1}b\} = \text{span}\{p_k(A)b \mid \text{degree}(p) < k\}$$

## Krylov subspace

the Krylov subspace of dimension  $k$  is

$$\mathcal{K}_k = \text{span}\{b, Ab, \dots, A^{k-1}b\} = \text{span}\{p_k(A)b \mid \text{degree}(p) < k\}$$

the iterates of the **Krylov sequence**  $x^{(1)}, x^{(2)}, \dots$ , minimize objective over successive Krylov subspaces

$$x^{(k)} = \underset{x \in \mathcal{K}_k}{\operatorname{argmin}} f(x) = \underset{x \in \mathcal{K}_k}{\operatorname{argmin}} \|Ax - b\| = \underset{x \in \mathcal{K}_k}{\operatorname{argmin}} \|x - x_\star\|_A$$

the CG algorithm generates the Krylov sequence

## Properties of Krylov sequence

- ▶  $f(x^{(k+1)}) \leq f(x^{(k)})$  (but  $\|r\|$  can increase)
- ▶  $x^{(n)} = x_*$
- ▶  $x^{(k)} = p_k(A)b$ , where  $p_k$  is a polynomial with degree  $< k$
- ▶ less obvious: there is a two-term recurrence

$$x^{(k+1)} = x^{(k)} + \alpha_k p^{(k)} \quad \text{where} \quad p^{(k)} = -r^{(k)} + \beta_k p^{(k-1)}$$

## Properties of Krylov sequence

- ▶  $f(x^{(k+1)}) \leq f(x^{(k)})$  (but  $\|r\|$  can increase)
- ▶  $x^{(n)} = x_*$
- ▶  $x^{(k)} = p_k(A)b$ , where  $p_k$  is a polynomial with degree  $< k$
- ▶ less obvious: there is a two-term recurrence

$$x^{(k+1)} = x^{(k)} + \alpha_k p^{(k)} \quad \text{where} \quad p^{(k)} = -r^{(k)} + \beta_k p^{(k-1)}$$

- ▶  $\alpha_k$  and  $\beta_k$  are determined by the CG algorithm



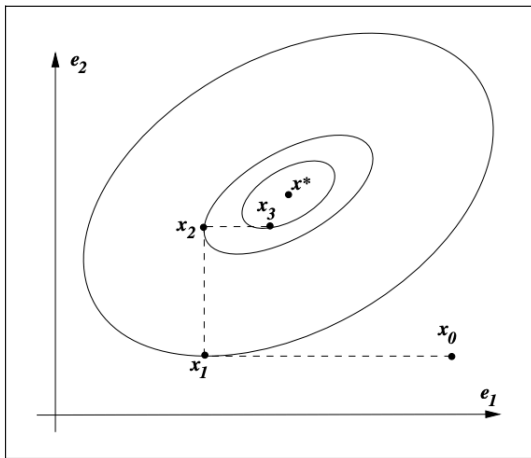
## Properties of Krylov sequence

- ▶  $f(x^{(k+1)}) \leq f(x^{(k)})$  (but  $\|r\|$  can increase)
- ▶  $x^{(n)} = x_*$
- ▶  $x^{(k)} = p_k(A)b$ , where  $p_k$  is a polynomial with degree  $< k$
- ▶ less obvious: there is a two-term recurrence

$$x^{(k+1)} = x^{(k)} + \alpha_k p^{(k)} \quad \text{where} \quad p^{(k)} = -r^{(k)} + \beta_k p^{(k-1)}$$

- ▶  $\alpha_k$  and  $\beta_k$  are determined by the CG algorithm
- ▶ can derive recurrence from optimality conditions:  
each new iterate  $x^{(k+1)}$  must have gradient (residual)  
orthogonal to  $\mathcal{K}_k$

## Coordinate descent does not solve in $n$ iterations



**Figure 5.2** Successive minimization along coordinate axes does not find the solution in  $n$  iterations, for a general convex quadratic.

## CG converges in $\text{Rank}(A)$ iterations

write (don't compute!) SVD of  $A = V\Lambda V^\top$  with

- ▶  $r = \text{Rank}(A)$
- ▶  $\Lambda \in \mathbf{R}^r \times r$  diagonal and positive
- ▶  $V \in \mathbf{R}^{n \times r}$ : orthonormal:  $V^\top V = I_r$

## CG converges in $\text{Rank}(A)$ iterations

write (don't compute!) SVD of  $A = V\Lambda V^\top$  with

- ▶  $r = \text{Rank}(A)$
- ▶  $\Lambda \in \mathbf{R}^r \times r$  diagonal and positive
- ▶  $V \in \mathbf{R}^{n \times r}$ : orthonormal:  $V^\top V = I_r$

characteristic polynomial of  $\Lambda$ :

$$\xi(s) = \det(sI_r - \Lambda) = (s - \lambda_1) \cdots (s - \lambda_r) = s^r + \alpha s^{r-1} + \cdots + \alpha_r$$

Cayley-Hamilton theorem

$$\begin{aligned}\xi(\Lambda) = 0 &= \Lambda^r + \alpha_1 \Lambda^{r-1} + \cdots + \alpha_r I_r \\ \Lambda^{-1} &= -(1/\alpha_r)(\Lambda^{r-1} + \alpha_1 \Lambda^{r-2} + \cdots + \alpha_{r-1} I_r)\end{aligned}$$

## CG converges in Rank( $A$ ) iterations

write (don't compute!) SVD of  $A = V\Lambda V^\top$  with

- ▶  $r = \text{Rank}(A)$
- ▶  $\Lambda \in \mathbf{R}^r \times r$  diagonal and positive
- ▶  $V \in \mathbf{R}^{n \times r}$ : orthonormal:  $V^\top V = I_r$

characteristic polynomial of  $\Lambda$ :

$$\xi(s) = \det(sI_r - \Lambda) = (s - \lambda_1) \cdots (s - \lambda_r) = s^r + \alpha s^{r-1} + \cdots + \alpha_r$$

Cayley-Hamilton theorem

$$\begin{aligned}\xi(\Lambda) = 0 &= \Lambda^r + \alpha_1 \Lambda^{r-1} + \cdots + \alpha_r I_r \\ \Lambda^{-1} &= -(1/\alpha_r)(\Lambda^{r-1} + \alpha_1 \Lambda^{r-2} + \cdots + \alpha_{r-1} I_r)\end{aligned}$$

write  $A^{-1} = V\Lambda^{-1}V^\top$  in terms of this decomposition:

$$\begin{aligned}A^{-1} = V\Lambda^{-1}V^\top &= -(1/\alpha_r)(V\Lambda^{r-1}V^\top + \alpha_1 V\Lambda^{r-2}V^\top + \cdots + \alpha_r I_r) \\ &= -(1/\alpha_r)(A^{r-1} + \alpha_1 A^{r-2} + \cdots + \alpha_{r-1} I)\end{aligned}$$

in particular,  $x_* = A^{-1}b \in \mathcal{K}_r$

# Outline

QR

Conjugate gradient

Preconditioned CG

Iterative refinement

## Matrix square root

$A \in \mathbf{S}_+^n$  has a square root  $A^{1/2} \in \mathbf{S}_+^n$ :

- ▶ if  $A = U\Lambda U^\top$  is the eigendecomposition of  $A$ ,
- ▶ then  $A^{1/2} = U\Lambda^{1/2}U^\top$

so  $A = A^{1/2}A^{1/2}$ .

## Preconditioning CG

for any  $P \succ 0$ ,

$$\begin{aligned} Ax = b & \iff P^{-1/2}Ax = P^{-1/2}b \\ & P^{-1/2}AP^{-1/2}z = P^{-1/2}b \end{aligned}$$

where  $x = P^{-1/2}z$ .



## Preconditioning CG

for any  $P \succ 0$ ,

$$\begin{aligned} Ax = b &\iff P^{-1/2}Ax = P^{-1/2}b \\ &\quad P^{-1/2}AP^{-1/2}z = P^{-1/2}b \end{aligned}$$

where  $x = P^{-1/2}z$ .

- preconditioning works well when  $\kappa(P^{-1/2}AP^{-1/2}) \ll \kappa(A)$

## Preconditioning CG

for any  $P \succ 0$ ,

$$\begin{aligned} Ax = b &\iff P^{-1/2}Ax = P^{-1/2}b \\ &\quad P^{-1/2}AP^{-1/2}z = P^{-1/2}b \end{aligned}$$

where  $x = P^{-1/2}z$ .

- ▶ preconditioning works well when  $\kappa(P^{-1/2}AP^{-1/2}) \ll \kappa(A)$

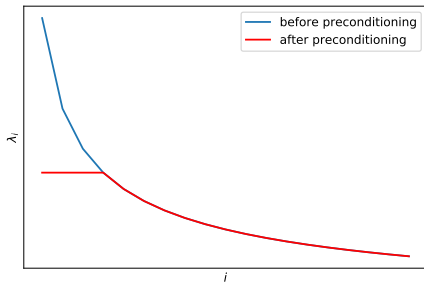
how to precondition?

- ▶ common heuristic: Jacobi preconditioning  $P = \mathbf{diag}(A)$
- ▶ incomplete Cholesky (best for structured sparsity)

## An optimal low-rank preconditioner

- ▶ suppose  $[A]_s = V_s \Lambda_s V_s^T$  is a best rank- $s$  apx to  $A \in \mathbf{S}_+^n$ .
- ▶ the best preconditioner using this information is

$$P_\star = \frac{1}{\lambda_{s+1}} V_s (\Lambda_s) V_s^T + (I - V_s V_s^T)$$



# Outline

QR

Conjugate gradient

Preconditioned CG

Iterative refinement

## Iterative refinement

want to solve  $Ax = b$ .

given approximate solution  $Ax^{(0)} \approx b$ , for  $k = 1, \dots$ ,

- ▶ compute residual  $r^{(k)} = b - Ax^{(k)}$
- ▶ use any method to solve  $A\delta^{(k)} = r^{(k)}$
- ▶  $x^{(k+1)} = x^{(k)} + \delta^{(k)}$